

# Matrix-Matrix Multiplication – A More Complex Kernel

- We can map every valid data element in a 2D array to a unique thread using threadIdx, blockIdx, blockDim, and gridDim variables:

```
// Calculate the row #
```

```
int Row = blockIdx.y*blockDim.y + threadIdx.y;
```


```
// Calculate the column #
```

```
int Col = blockIdx.x*blockDim.x + threadIdx.x;
```

- Matrix-Matrix multiplication between an  $I \times J$  matrix  $d\_M$  and a  $J \times K$  matrix  $d\_N$  produces an  $I \times K$  matrix  $d\_P$ .

# Matrix-Matrix Multiplication – A More Complex Kernel

- When performing a matrix-matrix multiplication, *each element of the product matrix  $d\_P$  is an inner product of a row of  $d\_M$  and a column of  $d\_N$* . The inner product between two vectors is the sum of products of corresponding elements. That is,  $d\_P_{\text{Row},\text{Col}} = \sum d\_M_{\text{Row},k} * d\_N_{k,\text{Col}}$ , for  $k = 0, 1, \dots, \text{Width}-1$ .



- We design a kernel where **each thread is responsible for calculating one  $d\_P$  element**.
- The  $d\_P$  element calculated by a thread is in row  **$\text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$**  and in column  **$\text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$** .


# Matrix-Matrix Multiplication – three variants

We will write program in CUDA to multiply **matrix-A** ( $h_a \times w_a$ ) and **matrix-B** ( $h_b \times w_b$ ) in 3 variations as follows:

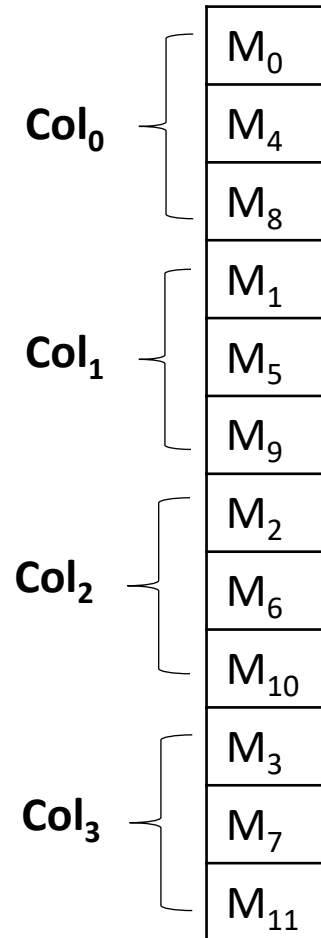
1. Each row of resultant matrix to be computed by one thread.
2. Each column of resultant matrix to be computed by one thread.
3. Each element of resultant matrix to be computed by one thread.

# Accessing Row & Col of a matrix

## Accessing Row



- Recall that a matrix  $M$  is linearized into an equivalent 1D array where the rows of  $M$  are placed one after another in the memory space, starting with the 0 row.
- For example, the beginning element of the 1 row is  $M[1 * \text{Width}]$  because we need to account for all elements of the 0 row.
- In general, the beginning element of the Row row is  $M[\text{Row} * \text{Width}]$ .
- Since all elements of a row are placed in consecutive locations, the  $k$  element of the Row row is at  $M[\text{Row} * \text{Width} + k]$ .



## Accessing Col

- The beginning element of the Col column is the Col element of the 0 row, which is  $M[\text{Col}]$ .
- Accessing each additional element in the Col column requires skipping over entire rows. This is because the next element of the same column is actually the same element in the next row.
- Therefore, the  $k$  element of the Col column is  $M[k * \text{Width} + \text{Col}]$ .

Matrix  $M$

M <sub>0</sub>	M <sub>1</sub>	M <sub>2</sub>	M <sub>3</sub>
M <sub>4</sub>	M <sub>5</sub>	M <sub>6</sub>	M <sub>7</sub>
M <sub>8</sub>	M <sub>9</sub>	M <sub>10</sub>	M <sub>11</sub>

# 1. Each row of resultant matrix to be computed by one thread

```
multiplyKernel_a<<<1,ha>>>(d_a, d_b, d_c, wa, wb);
```

```
__global__ void multiplyKernel_rowwise(int * a, int * b, int * c, int wa, int wb)
{
    int ridA = threadIdx.x;
    int sum;
    for(int cidB = 0; cidB < wb; cidB++)
    {
        sum= 0;
        for(k = 0; k< wa; k++)
        {
            sum += (a[ridA * wa + k] * b[k * wb + cidB]);
        }
        c[ridA * wb+ cidB] = sum;
    }
}
```

## 2. Each column of resultant matrix to be computed by one thread

```
multiplyKernel_b<<<1, wb>>>(d_a, d_b, d_c, ha, wa);
```

```
__global__ void multiplyKernel_colwise(int * a, int * b, int * c, int ha, int wa)
{
    int cidB = threadIdx.x;
    int wb = blockDim.x;
    int sum, k;
    for(ridA = 0; ridA < ha; ridA++)
    {
        sum = 0;
        for( k=0; k< wa; k++)
        {
            sum += (a[ridA * wa + k] * b[k * wb + cidB]);
        }
        c[ridA * wb + cidB] =sum;
    }
}
```

### 3. Each element of resultant matrix to be computed by one thread

```
multiplyKernel_b<<<(1, 1), (wb,ha)>>>(d_a, d_b, d_c, wa);
```

```
__global__ void multiplyKernel_elementwise(int * a, int * b, int * c, int wa)
{
    int ridA = threadIdx.y;
    int cidB= threadIdx.x;
    int wb = blockDim.x;
    int sum=0, k;

    for( k = 0; k < wa; k++)
    {
        sum += (a[ridA * wa + k] * b[k * wb + cidB]);
    }
    c[ridA * wb + cidB] =sum;
}
```

# Matrix-Matrix Multiplication—Kernel for thread-to-data mapping

- In the following kernel, we assume **square matrices** having dimension **Width\*Width**.
- Throughout the source code, instead of using a numerical value 16 for the block-width, the programmer can use the name BLOCK\_WIDTH by defining it. It helps in *autotuning*.

```
__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int Width)
{
    // Calculate the row index of the d_P element and d_M
    int Row = blockIdx.y*blockDim.y+threadIdx.y;
    // Calculate the column index of d_P and d_N
    int Col = blockIdx.x*blockDim.x+threadIdx.x;
    if ((Row < Width) && (Col < Width))
    {
        float Pvalue = 0;
        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += d_M[Row*Width+k]*d_N[k*Width+Col];

        d_P[Row*Width+Col] = Pvalue;
    }
}
```

Kernel code

```
#define BLOCK_WIDTH 16
```

```
// Setup the execution configuration
```

```
int NumBlocks = Width/BLOCK_WIDTH;
```

```
if (Width % BLOCK_WIDTH)
```

```
    NumBlocks++;
```

```
dim3 dimGrid(NumBlocks, NumBlocks);
```

```
dim3 dimBlock(BLOCK_WIDTH, BLOCK_WIDTH);
```

```
// Launch the device computation threads
```

```
matrixMulKernel<<dimGrid, dimBlock>>(M, N, P, Width);
```

Host code